

```
In [1]: import pandas as pd
import traceback
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: csv_movies_with_budget_df=pd.read_csv(r'./project_datasets/top-500-movies.csv', parse_dates=True)
csv_movies_with_budget_df.head()
```

Out[2]:

	rank	release_date	title	url	production_cost	domestic_gross	world
0	1	2019-04-23	Avengers: Endgame	/movie/Avengers-Endgame-(2019)#tab=summary	400000000	858373000	1214968618
1	2	2011-05-20	Pirates of the Caribbean: On Stranger Tides	/movie/Pirates-of-the-Caribbean-On-Stranger-Ti...	379000000	241071802	429496381
2	3	2015-04-22	Avengers: Age of Ultron	/movie/Avengers-Age-of-Ultron#tab=summary	365000000	459005868	678815482
3	4	2015-12-16	Star Wars Ep. VII: The Force Awakens	/movie/Star-Wars-Ep-VII-The-Force-Awakens#tab=...	306000000	936662225	1066483254
4	5	2018-04-25	Avengers: Infinity War	/movie/Avengers-Infinity-War#tab=summary	300000000	678815482	1066483254

Data Cleaning (Handling missing values and duplicate data: Removing unused information)

In this section, i will try to handle the problem of missing values, duplicated data and format:

- Remove the unused columns that are not useful in my analysis;
- Remove the movies having empty genres;
- Ensure there is no white spaces on the primay key, the movie title
- Remove duplicate rows
- Convert year of release as an int instead of float
- Changing format of release date into datetime format and budget_adj/revenue_adj format from float to int.

Step 1: Drop unused columns and duplicates or rows with null values

```
In [3]: #
csv_movies_with_budget_df.drop(['url', 'mpaa'], axis=1, inplace=True)
req_cols = ['title', 'genre']
csv_movies_with_budget_df[req_cols] = csv_movies_with_budget_df[req_cols].a
pplly(lambda col:col.str.strip())
```

```
In [4]: # drop duplicates
csv_movies_with_budget_df.drop_duplicates(inplace=True)
# confirm correction by rechecking for duplicates in the data
sum(csv_movies_with_budget_df.duplicated())
```

Out[4]: 0

```
In [5]: # drop rows with any null values
csv_movies_with_budget_df.dropna(subset=['genre'], inplace=True)
csv_movies_with_budget_df.dropna(inplace=True)
# checks if any of columns in the data have null values - should print False
csv_movies_with_budget_df.isnull().sum().any()
```

Out[5]: False

Step 2: Convert 'release_date' type from str to datetime

```
In [6]: # Convert 'release_date' type from str to datetime
csv_movies_with_budget_df['release_date']=pd.to_datetime(csv_movies_with_bu
dget_df['release_date'])
csv_movies_with_budget_df['release_date'].head()
```

```
Out[6]: 0    2019-04-23
1    2011-05-20
2    2015-04-22
3    2015-12-16
4    2018-04-25
Name: release_date, dtype: datetime64[ns]
```

Check for Missing values in the dataset

```
In [7]: csv_movies_with_budget_df.isnull().sum()
```

```
Out[7]: rank                0
release_date              0
title                    0
production_cost           0
domestic_gross            0
worldwide_gross           0
opening_weekend           0
genre                    0
theaters                 0
runtime                  0
year                    0
dtype: int64
```

All titles are available Although the release date column for a value is null, the year column is not null Some genre is null

```
In [8]: csv_movies_with_budget_df.isnull().any(axis=1).sum()
```

```
Out[8]: 0
```

```
In [9]: csv_movies_with_budget_df[csv_movies_with_budget_df['genre'].isnull()].year.value_counts()
```

```
Out[9]: Series([], Name: count, dtype: int64)
```

Only 1 movie belonging to 2015 has an empty genre, the other years are not relevant for my analysis.

Duplicates in the dataset

```
In [10]: sum(csv_movies_with_budget_df.duplicated())
```

```
Out[10]: 0
```

Step 3: Check Number of unique values in the dataset*

```
In [11]: # Returns the number of unique values in each column
csv_movies_with_budget_df.nunique()
```

```
Out[11]: rank                474
release_date              454
title                    472
production_cost           89
domestic_gross           474
worldwide_gross          474
opening_weekend          472
genre                     10
theaters                 404
runtime                   94
year                      30
dtype: int64
```

Step 4: Create Range Columns for all cost

```
In [12]: # Final number of movies
rows, col = csv_movies_with_budget_df.shape
print('After cleaning, we now have only {} entries of movies.'.format(row
s))
```

After cleaning, we now have only 474 entries of movies.

```
In [13]: #create the range function which takes 2 parameters, the divisor (convert t
o million) and the range
def create_ranges(df, col,div=1_000_000,step=5):
    cst_min,cst_max = df[col].agg(['min', 'max'])
    cost_range=[(start-step,start) for start in range(int(cst_min/div),int
(cst_max/div)+step,step)]
    labels=[f"{start}-{end}" for start, end in cost_range]
    # Define the bins
    bins=pd.IntervalIndex.from_tuples(cost_range)
    # Use "cut" method to group values into ranges
    return pd.cut(df[col].apply(lambda x : int(x/div)),
                    bins=bins, include_lowest=True
                    ).map(dict(zip(bins, labels)))
```

```
In [14]: #The idea here is to make the cost more readable. For example, a production
cost of 396,554,223 converts to a range of 396-401 million
csv_movies_with_budget_df['prod_cost_range_million'] =create_ranges(csv_movies_with_budget_df,"production_cost")
csv_movies_with_budget_df['worldwide_gross_range_million'] =create_ranges(csv_movies_with_budget_df,"worldwide_gross",step=10)
csv_movies_with_budget_df['domestic_gross_range_million'] =create_ranges(csv_movies_with_budget_df,"domestic_gross")
csv_movies_with_budget_df.head()
```

Out[14]:

	rank	release_date	title	production_cost	domestic_gross	worldwide_gross	opening
0	1	2019-04-23	Avengers: Endgame	400000000	858373000	2797800564	35
1	2	2011-05-20	Pirates of the Caribbean: On Stranger Tides	379000000	241071802	1045713802	9
2	3	2015-04-22	Avengers: Age of Ultron	365000000	459005868	1395316979	19
3	4	2015-12-16	Star Wars Ep. VII: The Force Awakens	306000000	936662225	2064615817	24
4	5	2018-04-25	Avengers: Infinity War	300000000	678815482	2048359754	25

```
In [15]: #Ensure all rows are accounted for
csv_movies_with_budget_df[csv_movies_with_budget_df["prod_cost_range_million"].isna()]
csv_movies_with_budget_df[csv_movies_with_budget_df["domestic_gross_range_million"].isna()]
csv_movies_with_budget_df[csv_movies_with_budget_df["worldwide_gross_range_million"].isna()]
```

Out[15]:

rank	release_date	title	production_cost	domestic_gross	worldwide_gross	opening_weekend
------	--------------	-------	-----------------	----------------	-----------------	-----------------

From the above, there are no rows with NAN

```
In [16]: # Count the occurrences of each range
csv_movies_with_budget_df.worldwide_gross_range_million.value_counts()
```

```
Out[16]: worldwide_gross_range_million
163-173      15
283-293      10
363-373      10
243-253      10
393-403      10
..
1793-1803     0
1803-1813     0
1813-1823     0
1823-1833     0
1463-1473     0
Name: count, Length: 290, dtype: int64
```

Step 5: Convert year column into a int datatype so year looks like 2019 instead of 2019.0

```
In [17]: csv_movies_with_budget_df["year"] = csv_movies_with_budget_df.year.astype
('int64')
csv_movies_with_budget_df
```

Out[17]:

	rank	release_date	title	production_cost	domestic_gross	worldwide_gross	opening_weekend_gross
0	1	2019-04-23	Avengers: Endgame	400000000	858373000	2797800564	101934000
1	2	2011-05-20	Pirates of the Caribbean: On Stranger Tides	379000000	241071802	1045713802	40000000
2	3	2015-04-22	Avengers: Age of Ultron	365000000	459005868	1395316979	68000000
3	4	2015-12-16	Star Wars Ep. VII: The Force Awakens	306000000	936662225	2064615817	50000000
4	5	2018-04-25	Avengers: Infinity War	300000000	678815482	2048359754	69000000
...
493	494	2008-02-14	The Spiderwick Chronicles	92500000	71195053	162839667	10000000
494	495	2004-10-22	The Incredibles	92000000	261441092	631441092	30000000
495	496	2013-02-06	A Good Day to Die Hard	92000000	67349198	304249198	10000000
496	497	2004-04-09	The Alamo	92000000	22406362	23911362	10000000
498	499	2013-12-19	The Secret Life of Walter Mitty	91000000	58236838	187861183	10000000

474 rows × 8 columns

Step 6: Convert key column title to lowercase

```
In [18]: csv_movies_with_budget_df['title'] = csv_movies_with_budget_df.title.str.lower()
```

```
In [19]: csv_movies_with_budget_df.reset_index(drop = True, inplace = True)
```

```
In [20]: csv_movies_with_budget_df.columns
```

```
Out[20]: Index(['rank', 'release_date', 'title', 'production_cost', 'domestic_gross',  
              'worldwide_gross', 'opening_weekend', 'genre', 'theaters', 'runtime',  
              'year', 'prod_cost_range_million', 'worldwide_gross_range_million',  
              'domestic_gross_range_million'],  
              dtype='object')
```

Step 7: Rearrange Columns

```
In [21]: re_order_cols = ['rank', 'title', 'year', 'release_date', 'genre', 'prod_cost',  
                        '_range_million', 'worldwide_gross_range_million',  
                        'domestic_gross_range_million', 'production_cost', 'domestic_gross',  
                        'worldwide_gross', 'opening_weekend', 'theaters', 'runtime']  
for i,col in enumerate(re_order_cols):  
    tmp = csv_movies_with_budget_df[col]  
    csv_movies_with_budget_df.drop(labels=[col],axis=1,inplace=True)  
    csv_movies_with_budget_df.insert(i,col,tmp)
```

```
In [22]: csv_movies_with_budget_df.head()
```

```
Out[22]:
```

	rank	title	year	release_date	genre	prod_cost_range_million	worldwide_gross_r
0	1	avengers: endgame	2019	2019-04-23	Action	396-401	
1	2	pirates of the caribbean: on stranger tides	2011	2011-05-20	Adventure	376-381	
2	3	avengers: age of ultron	2015	2015-04-22	Action	361-366	
3	4	star wars ep. vii: the force awakens	2015	2015-12-16	Adventure	301-306	
4	5	avengers: infinity war	2018	2018-04-25	Action	296-301	

Final data after cleanup


```
In [23]: csv_movies_with_budget_df.shape
```

```
Out[23]: (474, 14)
```

```
In [24]: csv_movies_with_budget_df.to_csv(r'./project_datasets/clean-500-movies.csv', index=False)
```

Are there any legal or regulatory guidelines for your data or project topic?

None as this is just a top 500 movies with their production cost.

Did you make any assumptions in cleaning/transforming the data?

I assumed the rank of the movie was opinioned and not based on any logic. I do not plan to use that as part of my analysis.

How was your data sourced / verified for credibility?

This comes from the top 500 movies set from kaggle.

Was your data acquired in an ethical way?

Yes, it was just the top 500 movies.